

# Handling the Positive-Definite Constraint in the Bayesian Learning Rule

Wu Lin (UBC)

Joint work with Mark Schmidt (UBC) and Mohammad Emtiyaz Khan (AIP, RIKEN)

# Bayesian Learning Rule

Bayes vs Optimization

Unified via Bayesian Learning Rule

A Natural-gradient Variational Inference Algorithm

**Challenges:** How to deal with constraints

**Contributions:** Improved Bayesian Learning Rule  
efficient updates for the positive-definite constraints  
using Riemannian Gradient Descent

# Two different worlds in ML

Bayesian:

**Given:** data

**Infer:** posterior

latent variables

model: prior + likelihood

Distinct methods in each world:

Inference methods

Unconstrained optimization:

**Given:** data

**Min:** objective

decision variables

model: regularizer + loss

v.s.

Optimizers

# Scope

## Challenges:

- unify methods in both worlds
- use inference techniques to design optimizers
- do optimization with uncertainty

## The Bayesian learning rule (Khan & Rue 2019):

- **variational inference** approach unifies methods
- help to design new optimizers
- give an uncertainty estimation for decision variables



**Issue of the rule:** expensive to deal with constraints

Example:

Gaussian posterior approximation (covariance matrix is **positive-definite**)

# Our contributions

- new (constrained) parametrization
- closed-form correction term to address the constraint issue
- Gaussian, mixture of Gaussians, gamma, and more approximations



# Outline

- Bayesian learning rule
- Key idea from a geometric viewpoint
- Example of our new rule
- Experimental results

# Variational inference

Variational Inference (VI): data  $\mathcal{D}$ , latent variables  $\mathbf{w}$ , approx:  $p(\mathbf{w}|\mathcal{D})$

$$\min_{\lambda} \mathcal{L}(\lambda) = \mathbb{E}_{q(\mathbf{w}|\lambda)} [\ell(\mathcal{D}, \mathbf{w})] + \mathbb{D}_{\text{KL}}[q(\mathbf{w}|\lambda) || \exp(-R(\mathbf{w}))]$$

negative ELBO

loss/negative log-likelihood

$$\ell(\mathcal{D}, \mathbf{w}) = -\log p(\mathcal{D}|\mathbf{w})$$

regularizer/negative log-prior

$$R(\mathbf{w}) = -\log p(\mathbf{w})$$


Example:

$$q(\mathbf{w}) = \mathcal{N}(\mathbf{w}|\mu, \mathbf{S}), \text{ Gaussian approximation } \lambda = \{\mu, \mathbf{S}\}$$

the mean

the precision matrix

**Issue: positive-definite constraint**



The VI framework bridges the gap between the Bayesian world and the optimization world.

# Handling the constraint

$\lambda'$  is a unconstrained transformation

Black-box gradient VI (BBVI):  $\lambda' = \lambda' - t \nabla_{\lambda'} \mathcal{L}(\lambda')$

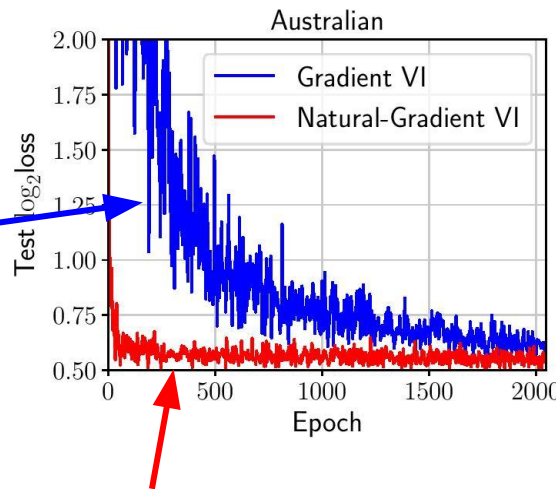
Natural-gradient VI (NGVI):  $\lambda' = \lambda' - t \mathbf{F}(\lambda')^{-1} \nabla_{\lambda'} \mathcal{L}(\lambda')$

Fisher matrix of  $q(\mathbf{w}|\lambda')$

Issue: **high** iteration cost in NGVI

Notation:  $\nabla$  denotes the ordinary/standard derivative in this talk

BBVI is **slow**



**Why NGVI instead of BBVI ?**

- **fast** iteration progress



# Bayesian learning rule

Bayesian learning rule is NGVI with low iteration cost

Under a right (**constrained**) parametrization  $\lambda$

$\nabla_m \mathcal{L}(\lambda) = \mathbf{F}(\lambda)^{-1} \mathcal{L}(\lambda)$ ,  $\mathbf{m}$  is an auxiliary parameter.

low iteration cost if easy to compute



Bayesian learning rule (BayesLRule):

$$\lambda = \lambda - t \nabla_m \mathcal{L}(\lambda)$$

line search for the step-size (**expensive**)



The rule:

- simple way to implement NGVI but gives rise to **constrained** optimization

# Example of the rule

Gaussian posterior approximation:

$$q(\mathbf{w}) = \mathcal{N}(\mathbf{w} | \mu, \mathbf{S})$$

the mean

the precision matrix

**Issue:**  $\mathbf{S}$  must be **positive-definite**

$\mathcal{L}(\lambda)$  can be re-written as  $\mathcal{L}(\lambda) = \mathbb{E}_{q(w|\lambda)}[\bar{\ell}(\mathbf{w})] + \mathbb{E}_{q(w|\lambda)}[\log q(\mathbf{w}|\lambda)]$

Bayesian learning rule for  $\min_{\lambda} \mathcal{L}(\lambda)$

$$\mathbf{S} = (1 - t)\mathbf{S} + t\mathbb{E}_q[\nabla_w^2 \bar{\ell}(\mathbf{w})]$$

$$\mu = \mu - t\mathbf{S}^{-1}\mathbb{E}_q[\nabla_w \bar{\ell}(\mathbf{w})]$$

The Bayesian learning rule

- bring the two worlds closer
- recover many optimizers (Khan & Rue 2019)

regularizer

loss

We define:  $\bar{\ell}(\mathbf{w}) := R(\mathbf{w}) + \ell(\mathcal{D}, \mathbf{w})$

Recall: minimize the negative ELBO  $\mathcal{L}(\lambda)$

**This work:** we address this constraint issue by using Riemannian gradient descent

Recall: Newton's update for  $\min_w \bar{\ell}(\mathbf{w})$

$$\mathbf{w} = \mathbf{w} - t[\nabla_w^2 \bar{\ell}(\mathbf{w})]^{-1} \nabla_w \bar{\ell}(\mathbf{w})$$

**require line search** for step-size  $t$



# Key idea of this work

Natural-gradient descent/Bayesian learning rule(BLR):

- like a “2nd-order” method in a Euclidean case
- 1st-order method in a (Riemannian) manifold
- 1st-order approximation of the geodesic at the starting point

## Challenges:


- hard to exactly compute a geodesic (manifold exponential map)
- NGD ignores the **curvature** of a geodesic

## This work:

- approximate a geodesic
- efficiently capture the curvature of a geodesic

natural-gradient descent/BLR

$$N(t) = x - t\mathbf{F}(\mathbf{x})^{-1}\nabla_x\mathcal{L}(\mathbf{x})$$



# Contributions of this work

Under a **new** (constrained) parametrization, our rule:

- efficiently approximate a geodesic as:  $\lambda = \lambda - t \nabla_m \mathcal{L}(\lambda)$  + correction  
- efficiently approximate **Christoffel symbols/Levi-Civita coefficients**
- efficiently compute natural-gradients  $\nabla_m \mathcal{L}(\lambda) = \mathbf{F}(\lambda)^{-1} \mathcal{L}(\lambda)$

capture the curvature of a geodesic

Example: Gaussian approximation  $q(\mathbf{w}) = \mathcal{N}(\mathbf{w} | \mu, \mathbf{S})$

the mean

the precision matrix  
(positive-definite constraint)

Bayesian learning rule (BayesLRule):

$$\begin{aligned}\mathbf{S} &= (1 - t)\mathbf{S} + t\mathbb{E}_q[\nabla_w^2 \bar{\ell}(\mathbf{w})] \\ \mu &= \mu - t\mathbf{S}^{-1}\mathbb{E}_q[\nabla_w \bar{\ell}(\mathbf{w})]\end{aligned}$$

require line search for step-size t

Our rule (iBayesLRule):

$$\begin{aligned}\mu &= \mu - t\mathbf{S}^{-1}\mathbb{E}_q[\nabla_w \bar{\ell}(\mathbf{w})] + \mathbf{0} \\ \mathbf{S} &= (1 - t)\mathbf{S} + t\mathbb{E}_q[\nabla_w^2 \bar{\ell}(\mathbf{w})] + \frac{t^2}{2}\mathbf{G}\mathbf{S}^{-1}\mathbf{G}\end{aligned}$$

no need to do line search

$$\mathbf{G} := \mathbf{S} - \mathbb{E}_q[\nabla_w^2 \bar{\ell}(\mathbf{w})]$$

same time complexity in terms of big O notation

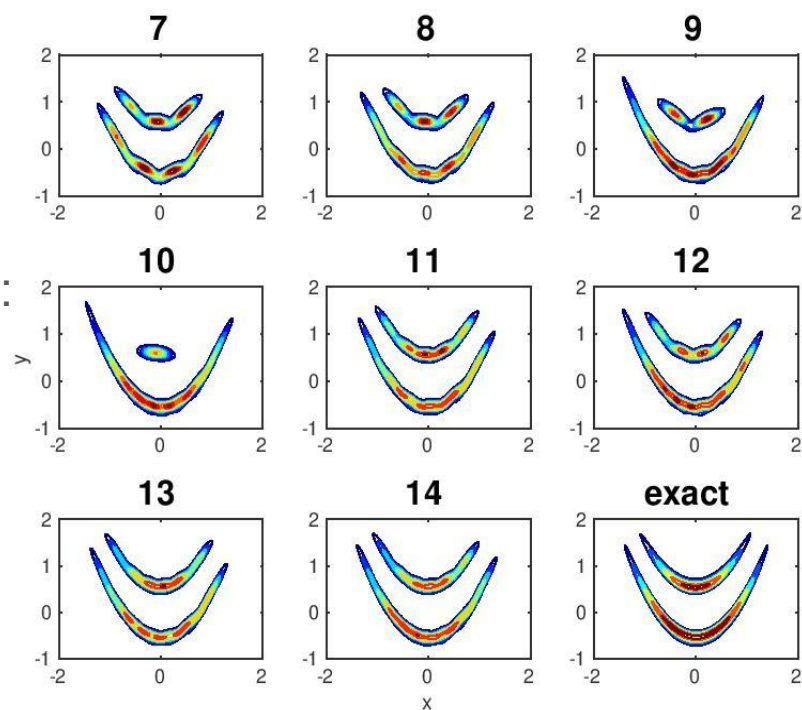
# More examples

Our approach can work on useful approximations:

- gamma
- inverse Gaussian
- Wishart

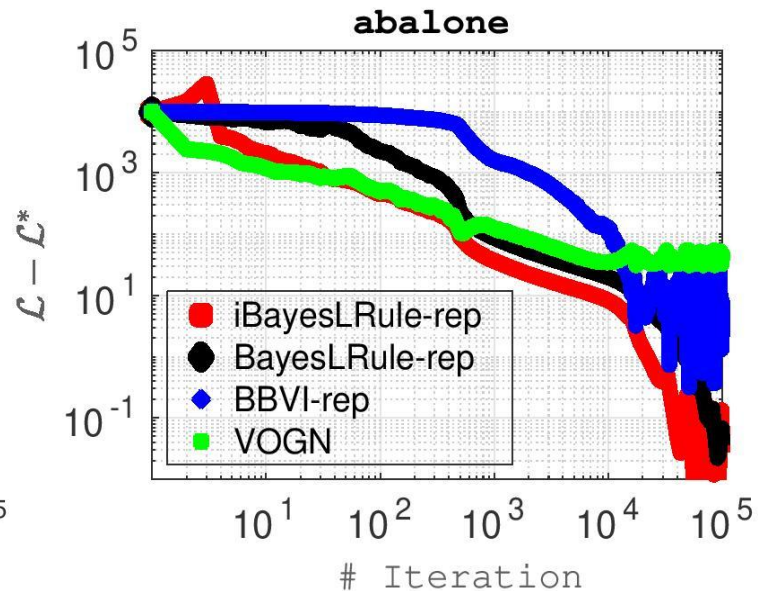
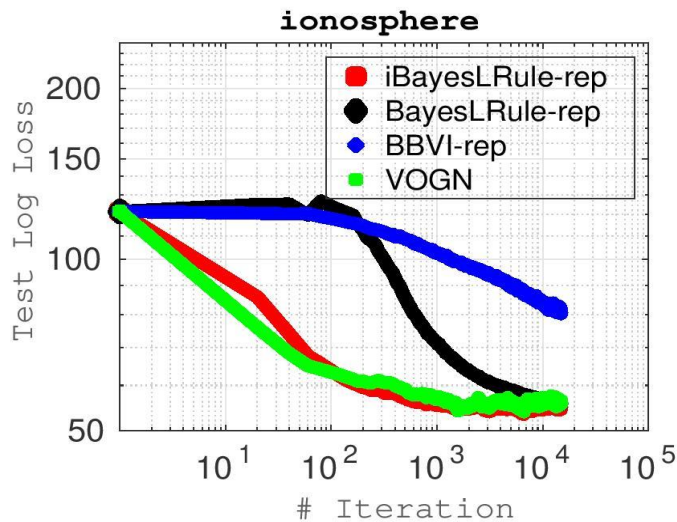
Mixture extensions:

- finite mixture models (e.g., mixture of Gaussians)
- skew Gaussian (continuous Gaussian mixture)
- Student's T (Gaussian scale mixtures)

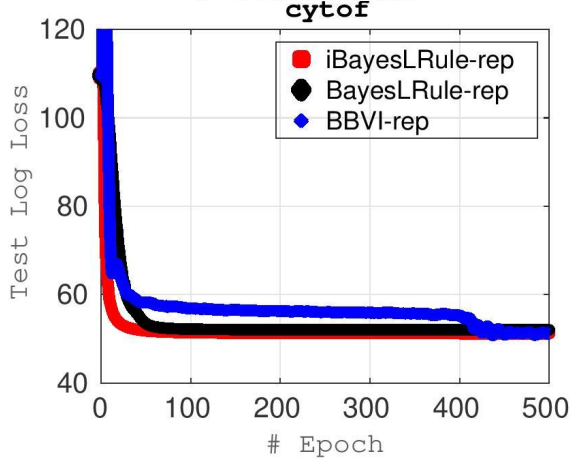


# Experimental results

Full Gaussian:



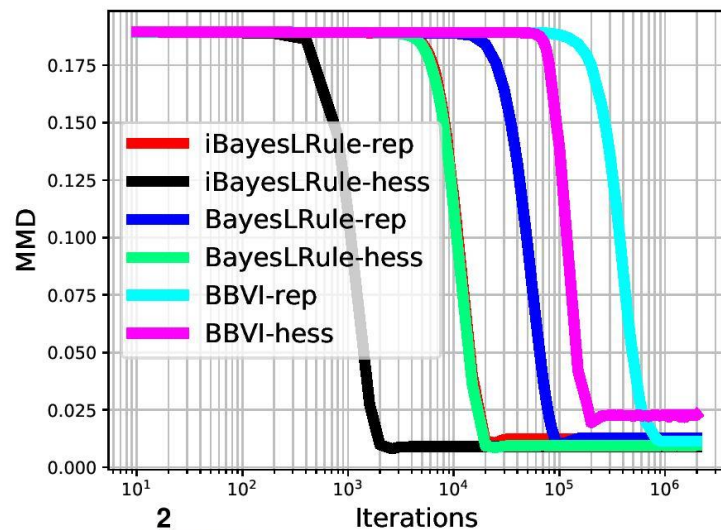
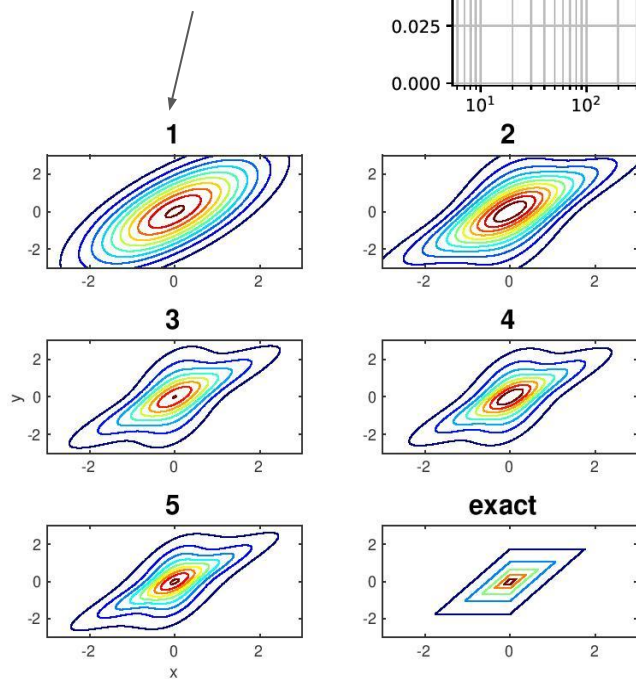
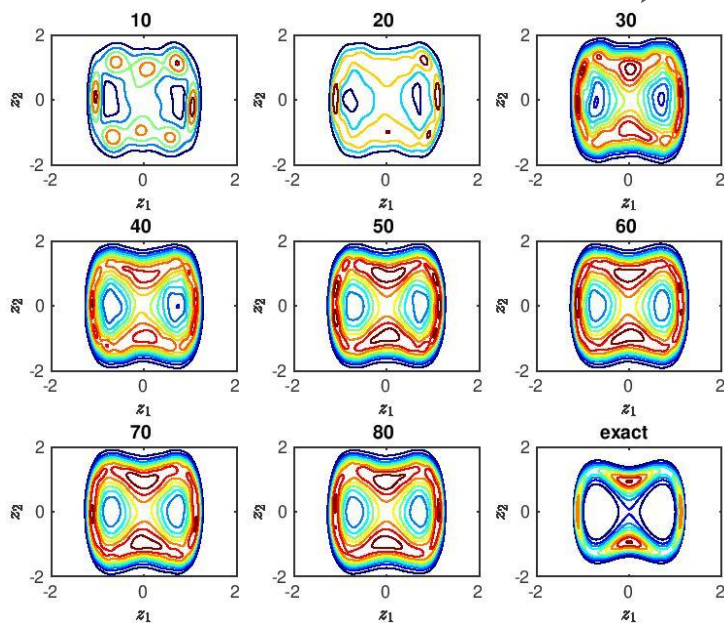
Gamma:



# More experimental results

Mixture of full Gaussians: 2-dim, 20-dim, 300-dim

# of mixing components





# Conclusion

**Improved Bayesian learning rule** (under a new constrained parametrization):

- handle positive-definite constraints
- use natural-gradients with low iteration cost
- no line search for many useful approximations
- as accurate as the original Bayesian learning rule
- easy to design new optimizers (e.g., ADAM-like optimizer)

**Thank You**

# Approaches to handle constraints

Example: Gaussian posterior approximation

**Issue:** parameter constraint (eg, covariance matrix)

Existing approaches to handle the constraint:

- block-box gradient VI: unconstrained, slow, accurate, **hard to design optimizers**
- Bayes learning rule: constrained, fast, accurate, **line search (expensive)**
- Gauss-Newton (VOGN): constrained, fast, specific approximation, **inaccurate**

**This work:**

constrained, fast, no line search, accurate, easy to design new optimizers

